

Yuezhou Hu

yuezhouhu@berkeley.edu | (341)333-8818 | yuezhouhu.github.io

Research Experience

-
- PhD Student**, Berkeley Artificial Intelligence Research Lab (BAIR), University of California, Berkeley, Advisor: [Prof. Kurt Keutzer](#) Sep. 2025 – Present
- Research Intern**, H. Milton Stewart School of Industrial and Systems Engineering, Georgia Institute of Technology, Advisor: [Prof. Tuo Zhao](#) Jul. 2024 – Aug. 2024
- Research Intern**, Statistical Artificial Intelligence & Learning Group, Tsinghua University, Advisor: [Prof. Jianfei Chen](#), [Prof. Jun Zhu](#) May 2022 – Jul. 2025

Education

-
- University of California, Berkeley**, Computer Science Ph.D Student Sep. 2025 – Present
- Department of Electrical Engineering and Computer Sciences
- Tsinghua University**, B.S. in Computer Science Sep. 2021 – Jul. 2025
- Department of Computer Science and Technology

Research Interests

My research interests primarily focus on **efficient machine learning**, particularly efficient training and inference. Recently, I am interested in:

- Diffusion Language Model; Speculative Decoding; Large Reasoning Model; Dynamic Sparse Training

Publications

-
- **AdaSPEC: Selective Knowledge Distillation for Efficient Speculative Decoders** [\[arXiv\]](#) [\[OpenReview\]](#)
*Yuezhou Hu**, Jiaxin Guo*, Xinyu Feng, Tuo Zhao
Neural Information Processing Systems (NeurIPS), 2025 (**Spotlight**)
 - **Pruning Large Language Models with Semi-Structural Adaptive Sparse Training** [\[arXiv\]](#) [\[Project page\]](#)
Weiyu Huang, *Yuezhou Hu*, Guohao Jian, Jun Zhu, Jianfei Chen
AAAI Conference on Artificial Intelligence (AAAI), 2025
 - **S-STE: Continuous Pruning Function for Efficient 2:4 Sparse Pre-training** [\[arXiv\]](#) [\[OpenReview\]](#) [\[Project page\]](#)
Yuezhou Hu, Jun Zhu, Jianfei Chen
Neural Information Processing Systems (NeurIPS), 2024
 - **Accelerating Transformer Pre-training with 2:4 Sparsity** [\[arXiv\]](#) [\[OpenReview\]](#) [\[PDF\]](#) [\[Project page\]](#)
Yuezhou Hu, Kang Zhao, Weiyu Huang, Jianfei Chen, Jun Zhu
International Conference on Machine Learning (ICML), 2024

Preprints

-
- **Arbitrage: Efficient Reasoning via Advantage-Aware Speculation** [\[arXiv\]](#) [\[Project page\]](#)
Monishwaran Maheswaran*, Rishabh Tiwari*, *Yuezhou Hu**, Kerem Dilmen, Coleman Hooper, Haocheng Xi, Nicholas Lee, Mehrdad Farajtabar, Michael Mahoney, Kurt Keutzer, Amir Gholami
 - **ParallelBench: Understanding the Trade-offs of Parallel Decoding in Diffusion LLMs** [\[arXiv\]](#)
Wonjun Kang, Kevin Galim, Seunghyuk Oh, Minjae Lee, Yuchen Zeng, Shuibai Zhang, Coleman Hooper, *Yuezhou Hu*, Hyung Il Koo, Nam Ik Cho, Kangwook Lee
 - **A Survey of Efficient Attention Methods: Hardware-efficient, Sparse, Compact, and Linear Attention**
Jintao Zhang, Rundong Su, Chunyu Liu, Jia Wei, Ziteng Wang, Pengl Zhang, Haoxu Wang, Huiqiang Jiang, Haofeng Huang, Chendong Xiang, Haocheng Xi, Shuo Yang, Xingyang Li, *Yuezhou Hu*, Tianyu Fu,

Tianchen Zhao, Yicheng Zhang, Boqun Cao, Youhe Jiang, Chang Chen, Kai Jiang, Huayu Chen, Min Zhao, Xiaoming Xu, Yi Wu, Fan Bao, Jun Zhu, Jianfei Chen

- **CAST: Continuous and Differentiable Semi-Structured Sparsity-Aware Training for Large Language Models** [[arXiv](#)]

Weiyu Huang, *Yuezhou Hu*, Jun Zhu, Jianfei Chen

- **Identifying Sensitive Weights via Post-quantization Integral** [[arXiv](#)]

Yuezhou Hu, Weiyu Huang, Zichen Liang, Chang Chen, Jintao Zhang, Jun Zhu, Jianfei Chen

Technical Skills

Deep learning programming: Python, Pytorch

GPU Programming: OpenAI Triton, C++ , CUDA

Others: Docker, Django, Rust

Language Skills: TOEFL 104 (R:28/L:27/S:23/W:26)

Honors

- | | |
|---|-----------|
| • Tsinghua Academic Preeminence Scholarship | Fall 2024 |
| • Tsinghua Innovation Future Scholarship | Fall 2024 |