

Yuezhou Hu

yuezhouhu@berkeley.edu | (341)333-8818 | yuezhouhu.github.io

Research Experience

PhD Student , Berkeley Artificial Intelligence Research Lab (BAIR), University of California, Berkeley, Advisor: Prof. Kurt Keutzer	Sep. 2025 – Present
Research Intern , H. Milton Stewart School of Industrial and Systems Engineering, Georgia Institute of Technology, Advisor: Prof. Tuo Zhao	Jul. 2024 – Aug. 2024
Research Intern , Statistical Artificial Intelligence & Learning Group, Tsinghua University, Advisor: Prof. Jianfei Chen , Prof. Jun Zhu	May 2022 – Jul. 2025

Education

University of California, Berkeley , Computer Science Ph.D Student Department of Electrical Engineering and Computer Sciences	Sep. 2025 – Present
Tsinghua University , B.S. in Computer Science Department of Computer Science and Technology	Sep. 2021 – Jul. 2025

Research Interests

My research interests primarily focus on **efficient machine learning**, particularly efficient training and inference. Recently, I am interested in:

- Diffusion Language Model; Speculative Decoding; Large Reasoning Model; Dynamic Sparse Training

Publications

3 first/cofirst-author papers, *equal contribution

- **AdaSPEC: Selective Knowledge Distillation for Efficient Speculative Decoders** [\[arXiv\]](#) [\[OpenReview\]](#)
Yuezhou Hu^{*}, Jiaxin Guo^{*}, Xinyu Feng, Tuo Zhao
Neural Information Processing Systems (NeurIPS), 2025 (**Spotlight**)
- **Pruning Large Language Models with Semi-Structural Adaptive Sparse Training** [\[arXiv\]](#) [\[Project page\]](#)
Weiyu Huang, *Yuezhou Hu*, Guohao Jian, Jun Zhu, Jianfei Chen
AAAI Conference on Artificial Intelligence (AAAI), 2025
- **S-STE: Continuous Pruning Function for Efficient 2:4 Sparse Pre-training** [\[arXiv\]](#) [\[OpenReview\]](#) [\[Project page\]](#)
Yuezhou Hu, Jun Zhu, Jianfei Chen
Neural Information Processing Systems (NeurIPS), 2024
- **Accelerating Transformer Pre-training with 2:4 Sparsity** [\[arXiv\]](#) [\[OpenReview\]](#) [\[PDF\]](#) [\[Project page\]](#)
Yuezhou Hu, Kang Zhao, Weiyu Huang, Jianfei Chen, Jun Zhu
International Conference on Machine Learning (ICML), 2024

Preprints

- **Arbitrage: Efficient Reasoning via Advantage-Aware Speculation** [\[arXiv\]](#) [\[Project page\]](#)
Monishwaran Maheswaran^{*}, Rishabh Tiwari^{*}, *Yuezhou Hu*^{*}, Kerem Dilmen, Coleman Hooper, Haocheng Xi, Nicholas Lee, Mehrdad Farajtabar, Michael Mahoney, Kurt Keutzer, Amir Gholami
- **ParallelBench: Understanding the Trade-offs of Parallel Decoding in Diffusion LLMs** [\[arXiv\]](#)
Wonjun Kang, Kevin Galim, Seunghyuk Oh, Minjae Lee, Yuchen Zeng, Shuibai Zhang, Coleman Hooper, *Yuezhou Hu*, Hyung Il Koo, Nam Ik Cho, Kangwook Lee
- **A Survey of Efficient Attention Methods: Hardware-efficient, Sparse, Compact, and Linear Attention**

Jintao Zhang, Rundong Su, Chunyu Liu, Jia Wei, Ziteng Wang, Pingle Zhang, Haoxu Wang, Huiqiang Jiang, Haofeng Huang, Chendong Xiang, Haocheng Xi, Shuo Yang, Xingyang Li, **Yuezhou Hu**, Tianyu Fu, Tianchen Zhao, Yicheng Zhang, Boqun Cao, Youhe Jiang, Chang Chen, Kai Jiang, Huayu Chen, Min Zhao, Xiaoming Xu, Yi Wu, Fan Bao, Jun Zhu, Jianfei Chen

- **CAST: Continuous and Differentiable Semi-Structured Sparsity-Aware Training for Large Language Models** [\[arXiv\]](#)

Weiyu Huang, **Yuezhou Hu**, Jun Zhu, Jianfei Chen

- **Identifying Sensitive Weights via Post-quantization Integral** [\[arXiv\]](#)

Yuezhou Hu, Weiyu Huang, Zichen Liang, Chang Chen, Jintao Zhang, Jun Zhu, Jianfei Chen

Technical Skills

Deep learning programming: Python, Pytorch

GPU Programming: OpenAI Triton, C++ , CUDA

Others: Docker, Django, Rust

Language Skills: TOEFL 104 (R:28/L:27/S:23/W:26)

Honors

- | | |
|---|-----------|
| • Tsinghua Academic Preeminence Scholarship | Fall 2024 |
| • 84 Future Innovation Scholarship | Fall 2024 |