

Yuezhou Hu

yuezhouhu@berkeley.edu | (341)333-8818 | yuezhouhu.github.io

Education

-
- Tsinghua University**, BS in Computer Science Sept 2021 – Jul. 2025
- Department of Computer Science and Technology
- Georgia Institute of Technology**, Visiting Student Jul. 2024 – Aug. 2024
- H. Milton Stewart School of Industrial and Systems Engineering
- University of California, Berkeley**, Computer Science PhD Student Sept 2025 – Present
- Department of Electrical Engineering and Computer Sciences

Research Interests

My research interests include **efficient machine learning**, particularly efficient training and inference. Recently, I am focusing on dynamic sparse training, quantization and speculative decoding. I want to break the monopoly of unicorns in AI and make deep learning affordable and accessible for every researcher.

Publications

-
- **Accelerating Transformer Pre-training with 2:4 Sparsity** [[arXiv](#)] [[OpenReview](#)] [[PDF](#)] [[Project page](#)]
Yuezhou Hu, Kang Zhao, Weiyu Huang, Jianfei Chen, Jun Zhu
International Conference on Machine Learning (ICML), 2024
 - **S-STE: Continuous Pruning Function for Efficient 2:4 Sparse Pre-training** [[arXiv](#)] [[OpenReview](#)] [[Project page](#)]
Yuezhou Hu, Jun Zhu, Jianfei Chen
Neural Information Processing Systems (NeurIPS), 2024
 - **Pruning Large Language Models with Semi-Structural Adaptive Sparse Training** [[arXiv](#)] [[Project page](#)]
Weiyu Huang, *Yuezhou Hu*, Guohao Jian, Jun Zhu, Jianfei Chen
AAAI Conference on Artificial Intelligence (AAAI), 2025
 - **AdaSPEC: Selective Knowledge Distillation for Efficient Speculative Decoders** [[arXiv](#)] [[OpenReview](#)]
Yuezhou Hu^{*}, Jiaxin Guo^{*}, Xinyu Feng, Tuo Zhao
Neural Information Processing Systems (NeurIPS), 2025 (Spotlight)

Preprints

-
- **Identifying Sensitive Weights via Post-quantization Integral** [[arXiv](#)]
Yuezhou Hu, Weiyu Huang, Zichen Liang, Chang Chen, Jintao Zhang, Jun Zhu, Jianfei Chen
 - **CAST: Continuous and Differentiable Semi-Structured Sparsity-Aware Training for Large Language Models** [[arXiv](#)]
Weiyu Huang, *Yuezhou Hu*, Jun Zhu, Jianfei Chen
 - **A Survey of Efficient Attention Methods: Hardware-efficient, Sparse, Compact, and Linear Attention**
Jintao Zhang, Rundong Su, Chunyu Liu, Jia Wei, Ziteng Wang, Pengle Zhang, Haoxu Wang, Huiqiang Jiang, Haofeng Huang, Chendong Xiang, Haocheng Xi, Shuo Yang, Xingyang Li, *Yuezhou Hu*, Tianyu Fu, Tianchen Zhao, Yicheng Zhang, Boqun Cao, Youhe Jiang, Chang Chen, Kai Jiang, Huayu Chen, Min Zhao, Xiaoming Xu, Yi Wu, Fan Bao, Jun Zhu, Jianfei Chen
 - **ParallelBench: Understanding the Trade-offs of Parallel Decoding in Diffusion LLMs** [[arXiv](#)]
Wonjun Kang, Kevin Galim, Seunghyuk Oh, Minjae Lee, Yuchen Zeng, Shuibai Zhang, Coleman Hooper, *Yuezhou Hu*, Hyung Il Koo, Nam Ik Cho, Kangwook Lee

Experience

Research Intern, Statistical Artificial Intelligence & Learning Group, Tsinghua University, Advisor: Prof. Jianfei Chen, Prof. Jun Zhu May 2022 – Present

Research Intern, H. Milton Stewart School of Industrial and Systems Engineering, Georgia Institute of Technology, Advisor: Prof. Tuo Zhao Jul. 2024 – Aug. 2024

PhD Student, Department of Electrical Engineering and Computer Sciences, University of California, Berkeley, Advisor: Prof. Kurt Keutzer Sept 2025 – Present

Technical Skills

Deep learning programming: Python, Pytorch

GPU Programming: OpenAI Triton, C++, CUDA

Others: Docker, Django, Rust

Language Skills: TOEFL 104 (R:28/L:27/S:23/W:26)

Honors

- Tsinghua Academic Preeminence Scholarship Fall 2024
5000 CNY (about 690 USD)
- Tsinghua Innovation Future Scholarship Fall 2024
25000 CNY (about 3450 USD) for a group of three students